

LoanPerformance_CaseStudy

January 6, 2019

1 Loan Performance Case Study

1.0.1 Imports

```
In [111]: import numpy as np
import copy
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import statsmodels.api as sm
import datetime as dt
import math
from dateutil.relativedelta import relativedelta
from scipy import stats
from sklearn.linear_model import LogisticRegression
from sklearn.linear_model import LogisticRegressionCV
from sklearn.preprocessing import OneHotEncoder
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.utils.fixes import signature
from sklearn import metrics
from collections import defaultdict
```

1.0.2 Read In Data

```
In [2]: df = pd.read_csv("loan.csv", usecols=["loan_amnt", "funded_amnt", "term", "int_rate", "grade",
                                              "annual_inc", "issue_d", "dti", "revol_bal", "total_pymnt", "total_pymnt_late", "churn"])
```

1.0.3 Data / Data Types

First we can take a quick look at our data.

```
In [3]: df.head()
```

```
Out[3]:
```

	loan_amnt	funded_amnt	term	int_rate	grade	annual_inc	issue_d	\
0	5000.0	5000.0	36 months	10.65	B	24000.0	Dec-2011	

1	2500.0	2500.0	60 months	15.27	C	30000.0	Dec-2011
2	2400.0	2400.0	36 months	15.96	C	12252.0	Dec-2011
3	10000.0	10000.0	36 months	13.49	C	49200.0	Dec-2011
4	3000.0	3000.0	60 months	12.69	B	80000.0	Dec-2011

	loan_status	dti	revol_bal	total_pymnt
0	Fully Paid	27.65	13648.0	5861.071414
1	Charged Off	1.00	1687.0	1008.710000
2	Fully Paid	8.72	2956.0	3003.653644
3	Fully Paid	20.00	5598.0	12226.302212
4	Current	17.94	27783.0	3242.170000

```
In [4]: df.shape
```

```
Out[4]: (887379, 11)
```

And the associated inferred datatypes from Pandas.

```
In [5]: df.dtypes
```

```
Out[5]: loan_amnt      float64
funded_amnt      float64
term              object
int_rate          float64
grade             object
annual_inc        float64
issue_d           object
loan_status       object
dti               float64
revol_bal         float64
total_pymnt       float64
dtype: object
```

Based on the data these inferred types seem reasonable, as we have float types inferred for apparent numeric columns and object (String) types inferred for the apparent string columns displayed above. One small change we will make, for ease of analysis, is to convert the 'issue_d' column to datetime.

```
In [6]: df["issue_d"] = pd.to_datetime(df["issue_d"])
```

```
In [7]: df.dtypes
```

```
Out[7]: loan_amnt      float64
funded_amnt      float64
term              object
int_rate          float64
grade             object
annual_inc        float64
issue_d           datetime64[ns]
```

```

loan_status      object
dti              float64
revol_bal        float64
total_pymnt      float64
dtype: object

```

1.1 1. Exploratory Data Analysis

1.1.1 Step 1 - Checking NULL / Empty Strings / Missing Values

First we will check numeric columns for NULL values.

```
In [8]: df.isna().sum()
```

```

Out[8]: loan_amnt      0
        funded_amnt    0
        term          0
        int_rate       0
        grade         0
        annual_inc     4
        issue_d        0
        loan_status     0
        dti            0
        revol_bal      0
        total_pymnt    0
        dtype: int64

```

Next we will check all values of object (String/str) columns for blank or odd values.

```
In [9]: df["term"].unique()
```

```
Out[9]: array([' 36 months', ' 60 months'], dtype=object)
```

```
In [10]: df["grade"].unique()
```

```
Out[10]: array(['B', 'C', 'A', 'E', 'F', 'D', 'G'], dtype=object)
```

```
In [11]: df["issue_d"].unique()
```

```

Out[11]: array(['2011-12-01T00:00:00.000000000', '2011-11-01T00:00:00.000000000',
                '2011-10-01T00:00:00.000000000', '2011-09-01T00:00:00.000000000',
                '2011-08-01T00:00:00.000000000', '2011-07-01T00:00:00.000000000',
                '2011-06-01T00:00:00.000000000', '2011-05-01T00:00:00.000000000',
                '2011-04-01T00:00:00.000000000', '2011-03-01T00:00:00.000000000',
                '2011-02-01T00:00:00.000000000', '2011-01-01T00:00:00.000000000',
                '2010-12-01T00:00:00.000000000', '2010-11-01T00:00:00.000000000',
                '2010-10-01T00:00:00.000000000', '2010-09-01T00:00:00.000000000',
                '2010-08-01T00:00:00.000000000', '2010-07-01T00:00:00.000000000',
                '2010-06-01T00:00:00.000000000', '2010-05-01T00:00:00.000000000',
                '2010-04-01T00:00:00.000000000', '2010-03-01T00:00:00.000000000',

```

```

'2010-02-01T00:00:00.000000000', '2010-01-01T00:00:00.000000000',
'2009-12-01T00:00:00.000000000', '2009-11-01T00:00:00.000000000',
'2009-10-01T00:00:00.000000000', '2009-09-01T00:00:00.000000000',
'2009-08-01T00:00:00.000000000', '2009-07-01T00:00:00.000000000',
'2009-06-01T00:00:00.000000000', '2009-05-01T00:00:00.000000000',
'2009-04-01T00:00:00.000000000', '2009-03-01T00:00:00.000000000',
'2009-02-01T00:00:00.000000000', '2009-01-01T00:00:00.000000000',
'2008-12-01T00:00:00.000000000', '2008-11-01T00:00:00.000000000',
'2008-10-01T00:00:00.000000000', '2008-09-01T00:00:00.000000000',
'2008-08-01T00:00:00.000000000', '2008-07-01T00:00:00.000000000',
'2008-06-01T00:00:00.000000000', '2008-05-01T00:00:00.000000000',
'2008-04-01T00:00:00.000000000', '2008-03-01T00:00:00.000000000',
'2008-02-01T00:00:00.000000000', '2008-01-01T00:00:00.000000000',
'2007-12-01T00:00:00.000000000', '2007-11-01T00:00:00.000000000',
'2007-10-01T00:00:00.000000000', '2007-09-01T00:00:00.000000000',
'2007-08-01T00:00:00.000000000', '2007-07-01T00:00:00.000000000',
'2007-06-01T00:00:00.000000000', '2013-12-01T00:00:00.000000000',
'2013-11-01T00:00:00.000000000', '2013-10-01T00:00:00.000000000',
'2013-09-01T00:00:00.000000000', '2013-08-01T00:00:00.000000000',
'2013-07-01T00:00:00.000000000', '2013-06-01T00:00:00.000000000',
'2013-05-01T00:00:00.000000000', '2013-04-01T00:00:00.000000000',
'2013-03-01T00:00:00.000000000', '2013-02-01T00:00:00.000000000',
'2013-01-01T00:00:00.000000000', '2012-12-01T00:00:00.000000000',
'2012-11-01T00:00:00.000000000', '2012-10-01T00:00:00.000000000',
'2012-09-01T00:00:00.000000000', '2012-08-01T00:00:00.000000000',
'2012-07-01T00:00:00.000000000', '2012-06-01T00:00:00.000000000',
'2012-05-01T00:00:00.000000000', '2012-04-01T00:00:00.000000000',
'2012-03-01T00:00:00.000000000', '2012-02-01T00:00:00.000000000',
'2012-01-01T00:00:00.000000000', '2014-12-01T00:00:00.000000000',
'2014-11-01T00:00:00.000000000', '2014-10-01T00:00:00.000000000',
'2014-09-01T00:00:00.000000000', '2014-08-01T00:00:00.000000000',
'2014-07-01T00:00:00.000000000', '2014-06-01T00:00:00.000000000',
'2014-05-01T00:00:00.000000000', '2014-04-01T00:00:00.000000000',
'2014-03-01T00:00:00.000000000', '2014-02-01T00:00:00.000000000',
'2014-01-01T00:00:00.000000000', '2015-12-01T00:00:00.000000000',
'2015-11-01T00:00:00.000000000', '2015-10-01T00:00:00.000000000',
'2015-09-01T00:00:00.000000000', '2015-08-01T00:00:00.000000000',
'2015-07-01T00:00:00.000000000', '2015-06-01T00:00:00.000000000',
'2015-05-01T00:00:00.000000000', '2015-04-01T00:00:00.000000000',
'2015-03-01T00:00:00.000000000', '2015-02-01T00:00:00.000000000',
'2015-01-01T00:00:00.000000000'], dtype='datetime64[ns]')

```

```
In [12]: df["loan_status"].unique()
```

```
Out[12]: array(['Fully Paid', 'Charged Off', 'Current', 'Default',
                'Late (31-120 days)', 'In Grace Period', 'Late (16-30 days)',
                'Does not meet the credit policy. Status:Fully Paid',
                'Does not meet the credit policy. Status:Charged Off', 'Issued'],
              dtype=object)
```

Based on the above analysis, it appears there are missing values in the 'annual_inc' column. All other columns appear to have no NULL, blank, or otherwise missing values. We can take a close look at the rows corresponding to NULL values in the 'annual_inc' column.

```
In [13]: df[df.isnull().any(axis=1)]
```

```
Out[13]:
```

	loan_amnt	funded_amnt	term	int_rate	grade	annual_inc	\
42449	5000.0	5000.0	36 months	7.43	A	NaN	
42450	7000.0	7000.0	36 months	7.75	A	NaN	
42480	6700.0	6700.0	36 months	7.75	A	NaN	
42533	6500.0	6500.0	36 months	8.38	A	NaN	

	issue_d	loan_status	dti	\
42449	2007-08-01	Does not meet the credit policy. Status:Fully ...	1.0	
42450	2007-08-01	Does not meet the credit policy. Status:Fully ...	1.0	
42480	2007-07-01	Does not meet the credit policy. Status:Fully ...	1.0	
42533	2007-06-01	Does not meet the credit policy. Status:Fully ...	4.0	

	revol_bal	total_pymnt
42449	0.0	5593.46
42450	0.0	7867.53
42480	0.0	7530.42
42533	0.0	7373.83

These NULL/NaN values are clustered in the mid-2007 period. This could also be due to internal issues with the LendingClub database or lack of requirements for reporting annual income within the database / front end. We will avoid dropping these records for now. If there is a business case to fill in these values then additional techniques such as imputation (using mean, MICE, or other methods) could be explored at a later time.

```
In [14]: df.shape
```

```
Out[14]: (887379, 11)
```

1.1.2 Step 2 - Initial EDA

```
In [15]: df_desc = df.describe()
display(df_desc)
```

	loan_amnt	funded_amnt	int_rate	annual_inc	\
count	887379.000000	887379.000000	887379.000000	8.873750e+05	
mean	14755.264605	14741.877625	13.246740	7.502759e+04	
std	8435.455601	8429.897657	4.381867	6.469830e+04	
min	500.000000	500.000000	5.320000	0.000000e+00	
25%	8000.000000	8000.000000	9.990000	4.500000e+04	
50%	13000.000000	13000.000000	12.990000	6.500000e+04	
75%	20000.000000	20000.000000	16.200000	9.000000e+04	
max	35000.000000	35000.000000	28.990000	9.500000e+06	

	dti	revol_bal	total_pymnt
count	887379.000000	8.873790e+05	887379.000000
mean	18.157039	1.692079e+04	7558.826684
std	17.190626	2.242679e+04	7871.243336
min	0.000000	0.000000e+00	0.000000
25%	11.910000	6.443000e+03	1914.590000
50%	17.650000	1.187500e+04	4894.999117
75%	23.950000	2.082900e+04	10616.814231
max	9999.000000	2.904836e+06	57777.579870

Initial Observations:

1. Loan application amount and funded amount appear to have very similar distributions, we will verify this later.
2. Interest rates seem have somewhat higher variance than expected.
3. Annual income seems to be right skewed as expected.
4. Debt-to-income seems to have relatively high variance.

In [16]: `df.corr()`


```
Out[16]:
```

	loan_amnt	funded_amnt	int_rate	annual_inc	dti	\
loan_amnt	1.000000	0.999263	0.145023	0.332698	0.020675	
funded_amnt	0.999263	1.000000	0.145160	0.332466	0.021075	
int_rate	0.145023	0.145160	1.000000	-0.072786	0.079903	
annual_inc	0.332698	0.332466	-0.072786	1.000000	-0.087410	
dti	0.020675	0.021075	0.079903	-0.087410	1.000000	
revol_bal	0.333580	0.333435	-0.035708	0.295784	0.067277	
total_pymnt	0.474626	0.473286	0.170506	0.160879	-0.041529	

	revol_bal	total_pymnt
loan_amnt	0.333580	0.474626
funded_amnt	0.333435	0.473286
int_rate	-0.035708	0.170506
annual_inc	0.295784	0.160879
dti	0.067277	-0.041529
revol_bal	1.000000	0.138328
total_pymnt	0.138328	1.000000

In [17]: `corr = df.corr()`
`sns.heatmap(corr,`
`xticklabels=corr.columns,`
`yticklabels=corr.columns, annot=True)`

Out[17]: <matplotlib.axes._subplots.AxesSubplot at 0x1a20751e10>



Correlation Observations:


1. Loan application amount and funding amount appear to be highly correlated (approaching 1.0).
2. Loan application amount, funding amount, annual income, revolving balance, and total payment all appear to be moderate-highly positively correlated.

All of these correlations seem reasonable and conform to my expectations about the loan application process, credit requirements, and nature of these fields.

Next we will plot distributions for the numeric columns.

```
In [18]: sns.distplot(df["loan_amnt"], kde=False, rug=True);
```

```
/Users/Samuel/anaconda3/lib/python3.6/site-packages/scipy/stats/stats.py:1713: FutureWarning: Us
return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval
```



```
In [19]: df.groupby(["loan_amnt"]).size().reset_index(name='count').sort_values(['count'], ascen
```


```
Out[19]:
```

	loan_amnt	count
371	10000.0	61837
451	12000.0	50183
571	15000.0	47210
771	20000.0	46932
1371	35000.0	36368
291	8000.0	27870
171	5000.0	27167
211	6000.0	26207
971	25000.0	24125
611	16000.0	23708
931	24000.0	22323
691	18000.0	22136
1171	30000.0	17588
1091	28000.0	14909
251	7000.0	14135
531	14000.0	13363
811	21000.0	12060
331	9000.0	11144
131	4000.0	10524
91	3000.0	9714

Loan application amount appears to be reasonably normally distribution in most cases. There are several values, including 10,000; 12,000; 15,000; 25,000; and 35,000 dollars, along with others


that appear to be overly represented based on standard distributional assumptions. This is likely due to human preference for these numbers.

```
In [20]: sns.distplot(df["funded_amnt"], kde=False, rug=True)
plt.show()
```



The distribution of funding amount closely mimics the distribution of loan application amount as expected.

```
In [21]: sns.distplot(df["int_rate"], kde=False, rug=True)
plt.show()
```




```
In [22]: df.groupby(["int_rate"]).size().reset_index(name='count').sort_values(['count'], ascending=False)
```

```
Out[22]:
```

	int_rate	count
110	10.99	34624
59	9.17	25720
287	15.61	25201
80	9.99	21553
38	7.89	20311


Interest rates appear to be reasonably distributed at first glance with no observable outliers. It is interesting to note the changes in density between different rate values in the plot.

```
In [23]: sns.distplot(df["annual_inc"].dropna(), kde=False, rug=True)
plt.show()
```



```
In [24]: thresh = 5*df_desc["annual_inc"]["std"]
          print(thresh)
          sns.distplot(df[abs(df["annual_inc"]-df_desc["annual_inc"]["mean"]) < thresh]["annual_i
          plt.show()
```

323491.50071200065




```
In [25]: df.groupby(["annual_inc"]).size().reset_index(name='count').sort_values(['count'], asce
```

```
Out[25]:
```

	annual_inc	count
24048	60000.0	34281
17805	50000.0	30575
26762	65000.0	25498
29378	70000.0	24121
11305	40000.0	23943


It appears there are some outliers for annual income and the distribution appears to be right skewed.

```
In [26]: sns.distplot(df["dti"], kde=False, rug=True)
plt.show()
```



```
In [27]: thresh = 5*df_desc["dti"]["std"]  
         print(thresh)  
         sns.distplot(df[abs(df["dti"]-df_desc["dti"]["mean"]) < thresh]["dti"], kde=False, rug=True,  
                       plt.show())
```

85.95312843970368




```
In [28]: df.groupby(["dti"]).size().reset_index(name='count').sort_values(['count'], ascending=False)
```

```
Out[28]:
```

	dti	count
1920	19.2	684
1440	14.4	674
1800	18.0	661
1320	13.2	638
1680	16.8	632


There appear to be some outliers for the 'debt-to-income' variable.

```
In [29]: sns.distplot(df["revol_bal"], kde=False, rug=True)
plt.show()
```



```
In [30]: thresh = 5*df_desc["revol_bal"]["std"]  
         print(thresh)  
         sns.distplot(df[abs(df["revol_bal"]-df_desc["revol_bal"]["mean"]) < thresh]["revol_bal"]  
         plt.show()
```

112133.95947981573




```
In [31]: df.groupby(["revol_bal"]).size().reset_index(name='count').sort_values(['count'], ascending=False)
```

```
Out[31]:
```

	revol_bal	count
0	0.0	3402
5235	5235.0	74
5466	5466.0	72
4479	4479.0	68
6969	6969.0	67

There appears to be some right-skew and outliers for the revolving balance variable.

```
In [32]: sns.distplot(df["total_pymnt"], kde=False, rug=True)
plt.show()
```

```
In [33]: df.groupby(["total_pymnt"]).size().reset_index(name='count').sort_values(['count'], asc
```

```
Out[33]:
```

	total_pymnt	count
0	0.00	17759
23309	648.57	135
35679	938.06	134
22207	623.86	131
61829	1566.46	117

There appears to be a large right skew for the total payment distribution.
Now we will explore our categorical variables.

```
In [34]: df["term"].value_counts()
```

```
Out[34]: 36 months    621125
        60 months    266254
        Name: term, dtype: int64
```

A 36 month loan term appears to be most common, representing 70.00% of the dataset.

```
In [35]: df["grade"].value_counts()
```

```
Out[35]: B    254535
        C    245860
        A    148202
```

```

D      139542
E       70705
F       23046
G        5489
Name: grade, dtype: int64

```

B grade loans are the most common, closely followed by C grade loans. The next most common grades, A and D grades, also have similar counts.

```
In [36]: df["issue_d"].value_counts()
```

```

Out[36]: 2015-10-01    48631
         2015-07-01    45962
         2015-12-01    44342
         2014-10-01    38782
         2015-11-01    37530
         2015-08-01    35886
         2015-04-01    35427
         2015-01-01    35107
         2015-05-01    31913
         2014-07-01    29306
         2015-09-01    28641
         2015-06-01    28485
         2015-03-01    25400
         2014-11-01    25054
         2015-02-01    23770
         2014-05-01    19099
         2014-04-01    19071
         2014-08-01    18814
         2014-06-01    17179
         2014-03-01    16513
         2014-01-01    15628
         2014-02-01    15269
         2013-12-01    15020
         2013-11-01    14676
         2013-10-01    14114
         2013-09-01    12987
         2013-08-01    12674
         2013-07-01    11910
         2013-06-01    10899
         2014-09-01    10606
         ...
         2009-12-01     658
         2009-10-01     604
         2009-09-01     507
         2009-08-01     446
         2009-07-01     411
         2009-06-01     406


```

2008-03-01	402
2009-05-01	359
2009-04-01	333
2009-03-01	324
2008-02-01	306
2008-01-01	305
2009-02-01	302
2009-01-01	269
2008-04-01	259
2008-12-01	253
2008-11-01	209
2007-12-01	172
2008-07-01	141
2008-06-01	124
2008-10-01	122
2008-05-01	115
2007-11-01	112
2007-10-01	105
2008-08-01	100
2007-08-01	74
2007-07-01	63
2008-09-01	57
2007-09-01	53
2007-06-01	24

Name: issue_d, Length: 103, dtype: int64

It appears that loan volumes have increased over time. We can look at this relationship below.

```
In [37]: count_df = df.groupby(["issue_d"]).size().reset_index(name="count").sort_values(["count"])
count_df[["issue_d", "count"]].set_index("issue_d").plot(legend=False)
plt.xlabel("Issue Date")
plt.ylabel("Total Loan Volume")
plt.show()
```



Loan volumes have indeed trended up over time. Interestingly, it appears that loan volumes have gotten much more variable during and after 2015.

```
In [38]: df["loan_status"].value_counts()
```

```
Out[38]: Current                601779
         Fully Paid             207723
         Charged Off            45248
         Late (31-120 days)      11591
         Issued                  8460
         In Grace Period         6253
         Late (16-30 days)       2357
         Does not meet the credit policy. Status:Fully Paid  1988
         Default                 1219
         Does not meet the credit policy. Status:Charged Off  761
         Name: loan_status, dtype: int64
```

The most common loan status is 'Current' followed by 'Fully Paid'. The other categories aside from 'Issued' represent loans which are at risk of default, are in default, or are charged off. 7.82% of all loans are in these categories.

```
In [39]: loan_status_agg = (df.groupby(["loan_status", "grade"]).size()/df.groupby(["grade"]).size())
         (loan_status_agg["Current"]+loan_status_agg["Fully Paid"])*100
```

```
Out[39]: grade
         A    96.490601
         B    93.614238
```

```

C    91.048971
D    87.431741
E    84.844070
F    79.471492
G    73.947896
dtype: float64

```

There appears to be a correlation between grades and associated loan status- 96.491% of A grade loans are “Current” or “Fully Paid” while only 73.948% of G grade loans are “Current” or “Fully Paid”.

```

In [40]: df["year"] = df["issue_d"].apply(lambda x: x.year)
df_grade_agg = df.groupby(["grade", "year"]).size()/df.groupby(["year"]).size()
df_grade_agg.head()

```

```

Out[40]: grade  year
A          2007    0.129353
          2008    0.132888
          2009    0.227798
          2010    0.225732
          2011    0.264905
dtype: float64

```

```

In [41]: df_year_agg = df.groupby(["year", "grade"]).size()/df.groupby(["year"]).size()
df_year_agg.head()

```

```


Out[41]: year  grade
2007  A          0.129353
      B          0.162521
      C          0.233831
      D          0.164179
      E          0.165837
dtype: float64

```


```

In [42]: df_grade_agg.unstack(level=0).plot(kind='bar', subplots=False)
plt.show()

```



```
In [43]: df_year_agg.unstack(level=0).plot(kind='bar', subplots=False)
plt.show()
```



Some interesting trends appear in the data: 1. We see that grade A loans peaked in popularity around 2011 (this may be due to tighter standards / more limited risk appetites after the financial crisis.) 2. Grade B loans peaked in popularity around 2012. 3. Grade C loans are now the most popular. 4. Grade F and G loans were more popular in 2007 than any time after.

1.1.3 Step 3 - Discussion of Outliers

Based on the above EDA there are apparent outliers in the following columns:

1. Annual Income
2. Debt to Income
3. Revolving Balance

We can initially quantify outliers using a 4 standard deviation cutoff from the mean, and will check for the presence of outliers in other columns.

```
In [44]: thresh1 = 4*df_desc["loan_amnt"]["std"]
         len(df[abs(df["loan_amnt"]-df_desc["loan_amnt"]["mean"]) >= thresh1])
```

```
Out[44]: 0
```

```
In [45]: thresh2 = 4*df_desc["funded_amnt"]["std"]
         len(df[abs(df["funded_amnt"]-df_desc["funded_amnt"]["mean"]) >= thresh2])
```

```
Out[45]: 0
```

```
In [46]: thresh3 = 4*df_desc["int_rate"]["std"]
         len(df[abs(df["int_rate"]-df_desc["int_rate"]["mean"]) >= thresh3])
```

```
Out[46]: 0
```

```
In [47]: thresh4 = 4*df_desc["annual_inc"]["std"]
         len(df[abs(df["annual_inc"]-df_desc["annual_inc"]["mean"]) >= thresh4])
```

```
Out[47]: 3641
```

```
In [48]: thresh5 = 4*df_desc["dti"]["std"]
         len(df[abs(df["dti"]-df_desc["dti"]["mean"]) >= thresh5])
```

```
Out[48]: 13
```

```
In [49]: thresh6 = 4*df_desc["revol_bal"]["std"]
         len(df[abs(df["revol_bal"]-df_desc["revol_bal"]["mean"]) >= thresh6])
```

```
Out[49]: 6893
```

```
In [50]: thresh7 = 4*df_desc["total_pymnt"]["std"]
         len(df[abs(df["total_pymnt"]-df_desc["total_pymnt"]["mean"]) >= thresh7])
```

```
Out [50]: 5526
```

Based on this rule, total payments appears to have outlier values as well.

```
In [51]: len(df[(abs(df["annual_inc"]-df_desc["annual_inc"]["mean"]) >= thresh4)) | (abs(df["dti"]
```

```
Out [51]: 15103
```

Given a somewhat crude rule of $[\text{abs}(\text{value}-\text{mean}) > 4*\text{sd}]$ representing an outlier for a given numeric attribute, we can identify 15,103 records corresponding to this rule when considering all numeric attributes. If there was time to further dive into working on outlier detection we could look at techniques such as one-class SVM or isolation forests. We will not exclude the identified outliers for now and will deal with them later during the modeling phase.

1.1.4 Step 4 - Summary of Findings

Note: The two(+) data visualizations and summary statistics to support the findings below can be found above.

Based on our exploratory data analysis, there are several important findings:

1. Four records with NULL or otherwise blank values were found.
2. Loan amount and funding amount appear to be highly correlated.
3. Loan amount, funding amount, annual incoming, revolving balance, and total payments all appear to be moderately-highly correlated.
4. Approximately 15,103 outliers were found in our numeric data using a 4 standard deviation cutoff.
5. Total loan volumes have trended up over time.
6. Loans that are in good standing make up the bulk of the loan portfolio (92.18%).
7. Higher grade loans appear to have better payback characteristics.
8. There appear to be time trends in the distribution of loan grades, potentially due to lender demand or regulatory requirements.

2. Business Analysis

We will subset to include only loans with 36 month terms.

```
In [52]: ba_df = copy.deepcopy(df)
         ba_df = ba_df[ba_df.term == " 36 months"]
```

```
In [53]: ba_df.head()
```

```
Out [53]:
```

	loan_amnt	funded_amnt	term	int_rate	grade	annual_inc	issue_d	\
0	5000.0	5000.0	36 months	10.65	B	24000.0	2011-12-01	
2	2400.0	2400.0	36 months	15.96	C	12252.0	2011-12-01	
3	10000.0	10000.0	36 months	13.49	C	49200.0	2011-12-01	
5	5000.0	5000.0	36 months	7.90	A	36000.0	2011-12-01	
7	3000.0	3000.0	36 months	18.64	E	48000.0	2011-12-01	

	loan_status	dti	revol_bal	total_pymnt	year
--	-------------	-----	-----------	-------------	------

0	Fully Paid	27.65	13648.0	5861.071414	2011
2	Fully Paid	8.72	2956.0	3003.653644	2011
3	Fully Paid	20.00	5598.0	12226.302212	2011
5	Fully Paid	11.20	7963.0	5631.377753	2011
7	Fully Paid	5.35	8221.0	3938.144334	2011

In [54]: ba_df.shape

Out[54]: (621125, 12)

Now we will exclude loans with less than 36 months of data. We can utilize a date difference from the maximum date (12/2015) to the loan issue date.

```
In [55]: end_date = max(ba_df["issue_d"])
         print(end_date)
```

2015-12-01 00:00:00

```
In [56]: ba_df["months_diff"] = ((end_date - ba_df["issue_d"])/np.timedelta64(1, 'M')).astype(int)
```

In [57]: ba_df.head()

```
Out[57]:
```

	loan_amnt	funded_amnt	term	int_rate	grade	annual_inc	issue_d	\
0	5000.0	5000.0	36 months	10.65	B	24000.0	2011-12-01	
2	2400.0	2400.0	36 months	15.96	C	12252.0	2011-12-01	
3	10000.0	10000.0	36 months	13.49	C	49200.0	2011-12-01	
5	5000.0	5000.0	36 months	7.90	A	36000.0	2011-12-01	
7	3000.0	3000.0	36 months	18.64	E	48000.0	2011-12-01	

	loan_status	dti	revol_bal	total_pymnt	year	months_diff
0	Fully Paid	27.65	13648.0	5861.071414	2011	48
2	Fully Paid	8.72	2956.0	3003.653644	2011	48
3	Fully Paid	20.00	5598.0	12226.302212	2011	48
5	Fully Paid	11.20	7963.0	5631.377753	2011	48
7	Fully Paid	5.35	8221.0	3938.144334	2011	48

```
In [58]: ba_df_valid = copy.deepcopy(ba_df[ba_df["months_diff"] >= 36])
```

In [59]: ba_df_valid.head()

```
Out[59]:
```

	loan_amnt	funded_amnt	term	int_rate	grade	annual_inc	issue_d	\
0	5000.0	5000.0	36 months	10.65	B	24000.0	2011-12-01	
2	2400.0	2400.0	36 months	15.96	C	12252.0	2011-12-01	
3	10000.0	10000.0	36 months	13.49	C	49200.0	2011-12-01	
5	5000.0	5000.0	36 months	7.90	A	36000.0	2011-12-01	
7	3000.0	3000.0	36 months	18.64	E	48000.0	2011-12-01	

	loan_status	dti	revol_bal	total_pymnt	year	months_diff
0	Fully Paid	27.65	13648.0	5861.071414	2011	48

2	Fully Paid	8.72	2956.0	3003.653644	2011	48
3	Fully Paid	20.00	5598.0	12226.302212	2011	48
5	Fully Paid	11.20	7963.0	5631.377753	2011	48
7	Fully Paid	5.35	8221.0	3938.144334	2011	48

```
In [60]: ba_df_valid.shape
```

```
Out[60]: (70200, 13)
```

2.0.1 Question 1

```
In [61]: def default_check(row):
          if row["loan_status"] == "Fully Paid":
              return 0
          else:
              return 1
```

```
ba_df_valid["default_flag"] = ba_df_valid.apply(default_check, axis=1)
```

```
In [62]: len(ba_df_valid[ba_df_valid["default_flag"] == 0]) / len(ba_df_valid)
```

```
Out[62]: 0.8432336182336182
```

84.323% of matured 36-month loans have been fully paid.

2.0.2 Question 2

```
In [63]: ba_df_valid["year"] = ba_df["issue_d"].apply(lambda x: x.year)
          ba_df_valid["month"] = ba_df["issue_d"].apply(lambda x: x.month)
```

```
# Setting model_df here to avoid duplicating transformations.
```

```
model_df = copy.deepcopy(ba_df_valid.drop(["months_diff", "loan_status", "term"], axis=
```

```
In [64]: grp_default_df = ba_df_valid[ba_df_valid["default_flag"] == 1].groupby(["year", "grade"
```

```
In [65]: grp_df = ba_df_valid.groupby(["year", "grade"]).size()
```

```
In [66]: grp_delta_df = grp_default_df / grp_df
          display(grp_delta_df)
```

year	grade	
2007	A	0.282051
	B	0.459184
	C	0.567376
	D	0.777778
	E	0.910000
	F	0.903846
	G	1.000000
2008	A	0.125786
	B	0.271044

	C	0.379310
	D	0.599045
	E	0.778947
	F	0.900901
	G	0.976744
2009	A	0.086451
	B	0.163322
	C	0.245549
	D	0.334149
	E	0.399351
	F	0.495238
	G	0.727273
2010	A	0.053370
	B	0.134403
	C	0.224155
	D	0.294493
	E	0.348214
	F	0.505495
	G	0.647059
2011	A	0.063990
	B	0.105676
	C	0.155243
	D	0.180809
	E	0.205882
	F	0.240741
	G	0.400000
2012	A	0.074203
	B	0.129363
	C	0.178265
	D	0.215827
	E	0.213333
	F	0.197674
	G	0.181818

dtype: float64

```
In [67]: grp_delta_df.idxmax()
```

```
Out[67]: (2007, 'G')
```

Grade 'G' loans issued in 2007 have the highest rate of default.

2.0.3 Question 3

```
In [68]: ba_df_valid["rate_of_return"] = (((ba_df_valid["total_pymnt"])/ba_df_valid["funded_amnt"]
```

```
In [69]: ror_df = ba_df_valid.groupby(["year", "grade"]).agg("mean")["rate_of_return"]
display(ror_df)
```

year	grade	
2007	A	0.027138
	B	-0.001888
	C	-0.008031
	D	-0.014353
	E	-0.015099
	F	-0.079863
	G	-0.044102
2008	A	0.023445
	B	-0.000428
	C	-0.004508
	D	-0.018365
	E	-0.003337
	F	-0.045138
	G	-0.007913
2009	A	0.021337
	B	0.014358
	C	0.011975
	D	0.013982
	E	-0.000270
	F	0.002615
	G	0.002985
2010	A	0.021611
	B	0.021892
	C	0.022227
	D	0.012545
	E	0.013981
	F	-0.007814
	G	-0.007685
2011	A	0.014910
	B	0.021848
	C	0.017563
	D	0.021422
	E	0.024065
	F	0.022267
	G	-0.020399
2012	A	0.014938
	B	0.022425
	C	0.021012
	D	0.022788
	E	0.033121
	F	0.031940
	G	0.038493

Name: rate_of_return, dtype: float64

In [70]: ror_df.idxmax()

Out[70]: (2012, 'G')

Interestingly, Grade 'G' loans issued in 2012 have the highest rate of return.

3. Logistic Regression Model

Note: For this question I will focus on 36-month loans using the criteria as defined above.

```
In [71]: # Using model_df as set above.
model_df.head()
```

```
Out[71]:
```

	loan_amnt	funded_amnt	int_rate	grade	annual_inc	issue_d	dti	\
0	5000.0	5000.0	10.65	B	24000.0	2011-12-01	27.65	
2	2400.0	2400.0	15.96	C	12252.0	2011-12-01	8.72	
3	10000.0	10000.0	13.49	C	49200.0	2011-12-01	20.00	
5	5000.0	5000.0	7.90	A	36000.0	2011-12-01	11.20	
7	3000.0	3000.0	18.64	E	48000.0	2011-12-01	5.35	

	revol_bal	total_pymnt	year	default_flag	month
0	13648.0	5861.071414	2011	0	12
2	2956.0	3003.653644	2011	0	12
3	5598.0	12226.302212	2011	0	12
5	7963.0	5631.377753	2011	0	12
7	8221.0	3938.144334	2011	0	12

```
In [72]: model_df.shape
```

```
Out[72]: (70200, 12)
```

3.0.1 Dealing with NULL/NaN Values and Outliers Identified Above

First we will drop the NULL / NaN values identified above in order to train our logistic regression model. This seems especially reasonable considering the limited number of observations (4). Additional imputation methods could be considered at a later time depending on business requirements. It is also a requirement for scikit's training and testing that no NULL/NaN values are present.

```
In [73]: model_df = model_df.dropna()
```

```
In [74]: model_df.head()
```

```
Out[74]:
```

	loan_amnt	funded_amnt	int_rate	grade	annual_inc	issue_d	dti	\
0	5000.0	5000.0	10.65	B	24000.0	2011-12-01	27.65	
2	2400.0	2400.0	15.96	C	12252.0	2011-12-01	8.72	
3	10000.0	10000.0	13.49	C	49200.0	2011-12-01	20.00	
5	5000.0	5000.0	7.90	A	36000.0	2011-12-01	11.20	
7	3000.0	3000.0	18.64	E	48000.0	2011-12-01	5.35	

	revol_bal	total_pymnt	year	default_flag	month
0	13648.0	5861.071414	2011	0	12
2	2956.0	3003.653644	2011	0	12

3	5598.0	12226.302212	2011	0	12
5	7963.0	5631.377753	2011	0	12
7	8221.0	3938.144334	2011	0	12

3.0.2 Relevant Column Selection

Loan application amount, funded amount, interest rate, grade, annual income, issue date, debt-to-income, and revolving balances are variables that should be present at loan issuance.

```
In [75]: model_df = model_df[["loan_amnt", "funded_amnt", "int_rate", "grade", "annual_inc", "is
```

```
In [76]: model_df.shape
```

```
Out[76]: (70196, 11)
```

3.0.3 Outlier Removal

We can also exclude records that were identified as outliers above. Additional analysis or methods such as imputation may be needed depending on business requirements / model output requirements regarding dataset completeness.

```
In [77]: model_df = model_df[((abs(model_df["annual_inc"]-df_desc["annual_inc"]["mean"] < thresh
```

```
In [78]: model_df.shape
```

```
Out[78]: (69636, 11)
```

3.0.4 Feature Selection

Since loan application amount is extremely highly correlated with funding amount and wouldn't seem to have a long term bearing on defaults we will exclude that column from consideration for our regression modeling.

```
In [79]: model_df = model_df[["funded_amnt", "int_rate", "grade", "annual_inc", "issue_d", "dti"
```

Because we are using L2 regularization we do not need to be as concerned about multicollinearity in this case. A deep dive into multicollinearity and some PCA / factor analysis would be a next step here and will help with regression coefficient interpretation.

Next we will need to one-hot encode our categorical variables in order to perform logistic regression.

```
In [80]: model_df_one_hot = pd.concat([model_df[["funded_amnt", "int_rate", "annual_inc", "revol
pd.get_dummies(model_df["grade"], prefix="grade"),
pd.get_dummies(model_df["year"], prefix="year")],axis=1)
```

```
In [81]: model_df_one_hot.corr()
```

Out [81]:

	funded_amnt	int_rate	annual_inc	revol_bal	dti	\
funded_amnt	1.000000	0.193897	0.411876	0.332515	0.051064	
int_rate	0.193897	1.000000	0.025042	0.093209	0.155424	
annual_inc	0.411876	0.025042	1.000000	0.394318	-0.171021	
revol_bal	0.332515	0.093209	0.394318	1.000000	0.238943	
dti	0.051064	0.155424	-0.171021	0.238943	1.000000	
default_flag	-0.027009	0.175632	-0.075310	-0.007097	0.037826	
grade_A	-0.093148	-0.771472	0.019019	-0.069182	-0.129491	
grade_B	-0.018733	-0.054535	-0.034928	0.015718	0.057536	
grade_C	-0.011272	0.359583	-0.025183	0.006638	0.052748	
grade_D	0.100831	0.491458	0.021496	0.029918	0.023116	
grade_E	0.105750	0.277830	0.048020	0.038999	-0.002711	
grade_F	0.047839	0.154211	0.027796	0.034636	0.007779	
grade_G	0.024373	0.119593	0.023640	0.031992	0.021610	
year_2007	-0.038301	-0.003857	-0.017830	-0.022314	-0.052788	
year_2008	-0.063426	0.004279	-0.018113	0.014716	-0.044663	
year_2009	-0.032857	0.036938	0.002861	0.011979	-0.094106	
year_2010	-0.056839	-0.071214	0.002612	-0.019971	-0.104124	
year_2011	-0.085263	-0.194544	0.007921	-0.049317	-0.094660	
year_2012	0.154434	0.184689	0.000150	0.045702	0.222112	

	default_flag	grade_A	grade_B	grade_C	grade_D	grade_E	\
funded_amnt	-0.027009	-0.093148	-0.018733	-0.011272	0.100831	0.105750	
int_rate	0.175632	-0.771472	-0.054535	0.359583	0.491458	0.277830	
annual_inc	-0.075310	0.019019	-0.034928	-0.025183	0.021496	0.048020	
revol_bal	-0.007097	-0.069182	0.015718	0.006638	0.029918	0.038999	
dti	0.037826	-0.129491	0.057536	0.052748	0.023116	-0.002711	
default_flag	1.000000	-0.145737	-0.048959	0.061076	0.103989	0.103938	
grade_A	-0.145737	1.000000	-0.456108	-0.324911	-0.226943	-0.105188	
grade_B	-0.048959	-0.456108	1.000000	-0.384847	-0.268807	-0.124592	
grade_C	0.061076	-0.324911	-0.384847	1.000000	-0.191486	-0.088754	
grade_D	0.103989	-0.226943	-0.268807	-0.191486	1.000000	-0.061993	
grade_E	0.103938	-0.105188	-0.124592	-0.088754	-0.061993	1.000000	
grade_F	0.089157	-0.051861	-0.061428	-0.043758	-0.030564	-0.014166	
grade_G	0.096722	-0.035330	-0.041848	-0.029811	-0.020822	-0.009651	
year_2007	0.125294	-0.031406	-0.035488	0.005550	0.013088	0.076616	
year_2008	0.147329	-0.059744	-0.039969	0.012284	0.033379	0.104570	
year_2009	0.042324	-0.031797	-0.045893	0.028630	0.032054	0.052127	
year_2010	0.007854	0.003678	-0.035975	0.009628	0.021977	0.019603	
year_2011	-0.066537	0.131843	-0.016968	-0.071865	-0.044172	-0.026590	
year_2012	-0.050162	-0.064681	0.083132	0.030890	-0.010602	-0.071117	

	grade_F	grade_G	year_2007	year_2008	year_2009	year_2010	\
funded_amnt	0.047839	0.024373	-0.038301	-0.063426	-0.032857	-0.056839	
int_rate	0.154211	0.119593	-0.003857	0.004279	0.036938	-0.071214	
annual_inc	0.027796	0.023640	-0.017830	-0.018113	0.002861	0.002612	
revol_bal	0.034636	0.031992	-0.022314	0.014716	0.011979	-0.019971	
dti	0.007779	0.021610	-0.052788	-0.044663	-0.094106	-0.104124	

default_flag	0.089157	0.096722	0.125294	0.147329	0.042324	0.007854
grade_A	-0.051861	-0.035330	-0.031406	-0.059744	-0.031797	0.003678
grade_B	-0.061428	-0.041848	-0.035488	-0.039969	-0.045893	-0.035975
grade_C	-0.043758	-0.029811	0.005550	0.012284	0.028630	0.009628
grade_D	-0.030564	-0.020822	0.013088	0.033379	0.032054	0.021977
grade_E	-0.014166	-0.009651	0.076616	0.104570	0.052127	0.019603
grade_F	1.000000	-0.004758	0.085718	0.086391	0.043695	0.012521
grade_G	-0.004758	1.000000	0.083960	0.099110	0.032167	0.003615
year_2007	0.085718	0.083960	1.000000	-0.017044	-0.025917	-0.035360
year_2008	0.086391	0.099110	-0.017044	1.000000	-0.052686	-0.071883
year_2009	0.043695	0.032167	-0.025917	-0.052686	1.000000	-0.109303
year_2010	0.012521	0.003615	-0.035360	-0.071883	-0.109303	1.000000
year_2011	-0.018635	-0.022261	-0.045953	-0.093417	-0.142047	-0.193803
year_2012	-0.063385	-0.052646	-0.101859	-0.207065	-0.314859	-0.429578

	year_2011	year_2012
funded_amnt	-0.085263	0.154434
int_rate	-0.194544	0.184689
annual_inc	0.007921	0.000150
revol_bal	-0.049317	0.045702
dti	-0.094660	0.222112
default_flag	-0.066537	-0.050162
grade_A	0.131843	-0.064681
grade_B	-0.016968	0.083132
grade_C	-0.071865	0.030890
grade_D	-0.044172	-0.010602
grade_E	-0.026590	-0.071117
grade_F	-0.018635	-0.063385
grade_G	-0.022261	-0.052646
year_2007	-0.045953	-0.101859
year_2008	-0.093417	-0.207065
year_2009	-0.142047	-0.314859
year_2010	-0.193803	-0.429578
year_2011	1.000000	-0.558268
year_2012	-0.558268	1.000000

```
In [82]: predictors = model_df_one_hot.drop("default_flag", axis=1)
         target = model_df_one_hot["default_flag"]
```

```
In [83]: predictors.head()
```

```
Out[83]:
```

	funded_amnt	int_rate	annual_inc	revol_bal	dti	grade_A	grade_B	\
0	5000.0	10.65	24000.0	13648.0	27.65	0	1	
2	2400.0	15.96	12252.0	2956.0	8.72	0	0	
3	10000.0	13.49	49200.0	5598.0	20.00	0	0	
5	5000.0	7.90	36000.0	7963.0	11.20	1	0	
7	3000.0	18.64	48000.0	8221.0	5.35	0	0	

	grade_C	grade_D	grade_E	grade_F	grade_G	year_2007	year_2008	\
0	0	0	0	0	0	0	0	
2	1	0	0	0	0	0	0	
3	1	0	0	0	0	0	0	
5	0	0	0	0	0	0	0	
7	0	0	1	0	0	0	0	

	year_2009	year_2010	year_2011	year_2012
0	0	0	1	0
2	0	0	1	0
3	0	0	1	0
5	0	0	1	0
7	0	0	1	0

```
In [84]: target.head()
```

```
Out[84]: 0    0
          2    0
          3    0
          5    0
          7    0
          Name: default_flag, dtype: int64
```

3.0.5 Scaling

We now need to scale our numerical predictors to be mean 0 and unit variance. This is especially critical due to the regularization incorporated in the sklearn logistic regression method and is generally best practice for preprocessing.

```
In [85]: scalar = StandardScaler()
          scaled_predictors = copy.deepcopy(predictors)
          scaled_predictors[["funded_amnt", "int_rate", "annual_inc", "dti", "revol_bal"]] = scalar.fit_transform(predictors[["funded_amnt", "int_rate", "annual_inc", "dti", "revol_bal"]])
```

```
In [86]: scaled_predictors.describe()
```

```
Out[86]:
```

	funded_amnt	int_rate	annual_inc	revol_bal	dti	\
count	6.963600e+04	6.963600e+04	6.963600e+04	6.963600e+04	6.963600e+04	
mean	5.685121e-15	6.983701e-15	-8.190677e-17	-6.309718e-16	1.865279e-17	
std	1.000007e+00	1.000007e+00	1.000007e+00	1.000007e+00	1.000007e+00	
min	-1.497521e+00	-1.805667e+00	-1.685277e+00	-1.020973e+00	-2.011219e+00	
25%	-7.326330e-01	-8.446114e-01	-6.590923e-01	-6.381062e-01	-7.582490e-01	
50%	-1.439690e-01	4.464156e-02	-2.416595e-01	-2.555928e-01	-1.969854e-02	
75%	5.271830e-01	6.549674e-01	4.181537e-01	3.111469e-01	7.256029e-01	
max	3.676723e+00	3.571275e+00	7.231735e+00	7.347348e+00	2.713073e+00	

	grade_A	grade_B	grade_C	grade_D	grade_E	\
count	69636.000000	69636.000000	69636.000000	69636.000000	69636.000000	
mean	0.278017	0.350752	0.215162	0.117971	0.027931	
std	0.448025	0.477209	0.410937	0.322576	0.164776	

min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	0.000000	0.000000
50%	0.000000	0.000000	0.000000	0.000000	0.000000
75%	1.000000	1.000000	0.000000	0.000000	0.000000
max	1.000000	1.000000	1.000000	1.000000	1.000000

	grade_F	grade_G	year_2007	year_2008	year_2009 \
count	69636.000000	69636.000000	69636.000000	69636.000000	69636.000000
mean	0.006936	0.003231	0.008315	0.033488	0.074171
std	0.082994	0.056751	0.090806	0.179909	0.262052
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	0.000000	0.000000
50%	0.000000	0.000000	0.000000	0.000000	0.000000
75%	0.000000	0.000000	0.000000	0.000000	0.000000
max	1.000000	1.000000	1.000000	1.000000	1.000000

	year_2010	year_2011	year_2012
count	69636.000000	69636.000000	69636.000000
mean	0.129775	0.201189	0.553062
std	0.336058	0.400892	0.497180
min	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000
50%	0.000000	0.000000	1.000000
75%	0.000000	0.000000	1.000000
max	1.000000	1.000000	1.000000

3.0.6 Modeling Training and Testing

We will now split our dataset into a train and test set. This is critical to identify how our regression performs out of sample.

```
In [87]: predictors_train, predictors_test, target_train, target_test = train_test_split(scaled_
```

Next we can get our class probabilities of our outcome variables for our train and test sets.

```
In [88]: target_train.value_counts()/len(target_train)
```

```
Out [88]: 0    0.846236
          1    0.153764
          Name: default_flag, dtype: float64
```

```
In [89]: target_test.value_counts()/len(target_test)
```

```
Out [89]: 0    0.84356
          1    0.15644
          Name: default_flag, dtype: float64
```

Next we can build and fit our regression model. A few notes on the fitting:

1. We will use CV to help with hyperparameter selection for regularization.

2. I have found the regularization selection (L1/L2) does not change the resultant coefficients in this case. In general L2 regularization is preferred, especially in cases of potential multicollinearity, so that is my selection.
3. We will use the balanced class weight parameter which should help in this unbalanced case. Testing showed that removing this parameter increases precision for both classes but drastically reduces recall. If we considered recall as being more important than precision for this case, i.e. a false negative is more damaging than a false positive for the loan default scenario, then this is valuable.

```
In [90]: clf = LogisticRegressionCV(cv=5, random_state=0, solver='lbfgs', penalty='l2', class_weight='balanced')
```

```
In [91]: pd.set_option('display.float_format', lambda x: '%.3f' % x)
         coefficients = pd.concat([pd.DataFrame(predictors.columns),pd.DataFrame(np.transpose(np.array([
         display(coefficients)
```

```

           0      0
0  funded_amnt 0.958
1      int_rate 1.144
2  annual_inc  0.775
3   revol_bal  1.014
4         dti  1.073
5    grade_A  0.303
6    grade_B  0.501
7    grade_C  0.724
8    grade_D  0.883
9    grade_E  1.230
10   grade_F  2.052
11   grade_G  4.826
12  year_2007  4.320
13  year_2008  2.090
14  year_2009  0.810
15  year_2010  0.630
16  year_2011  0.475
17  year_2012  0.541
```

```
In [92]: clf.score(predictors_train, target_train)
```

```
Out[92]: 0.6509130658436214
```

```
In [93]: probabilities_train = clf.predict_proba(predictors_train)[::,1]
         print(probabilities_train)
```

```
[0.44438477 0.43466474 0.28810599 ... 0.68830438 0.4625606  0.50112447]
```

```
In [94]: predicted_train = clf.predict(predictors_train)
         print(predicted_train)
```

```
[0 0 0 ... 1 0 1]
```

```
In [95]: predicted_train_df = pd.DataFrame({'Probabilities':probabilities_train, 'Predicted':pre
        predicted_train_df.head()
```

```
Out[95]:
```

	Probabilities	Predicted
0	0.444	0
1	0.435	0
2	0.288	0
3	0.376	0
4	0.480	0

```
In [96]: clf.score(predictors_test, target_test)
```

```
Out[96]: 0.6458659704090514
```

```
In [97]: probabilities_test = clf.predict_proba(predictors_test)[:,:1]
        print(probabilities_test)
```

```
[0.3042541  0.35938146 0.43413551 ... 0.49756599 0.91573313 0.5690408 ]
```

```
In [98]: predicted_test = clf.predict(predictors_test)
        print(predicted_test)
```

```
[0 0 0 ... 0 1 1]
```

```
In [99]: np.bincount(predicted_test)
```

```
Out[99]: array([14043,  8937])
```

```
In [100]: predicted_test_df = pd.DataFrame({'Probabilities': probabilities_test, 'Predicted': pr
        predicted_test_df.head()
```

```
Out[100]:
```

	Probabilities	Predicted
0	0.304	0
1	0.359	0
2	0.434	0
3	0.644	1
4	0.292	0

```
In [101]: metrics.confusion_matrix(target_test, predicted_test)
```

```
Out[101]: array([[12645,  6740],
                [ 1398,  2197]])
```

```
In [102]: tn, fp, fn, tp = metrics.confusion_matrix(target_test, predicted_test).ravel()
```

```
In [103]: print(tn, fp, fn, tp)
```

12645 6740 1398 2197

```
In [104]: print(metrics.classification_report(target_test, predicted_test))
```


	precision	recall	f1-score	support
0	0.90	0.65	0.76	19385
1	0.25	0.61	0.35	3595
avg / total	0.80	0.65	0.69	22980

Thus our model yields a precision of our positive class of $2178/(2178+6662)$ or .246 and a recall rate of $2,178/(2,178+1,430)$ or .604 for the given decision threshold (.5).

```
In [105]: auc_score = metrics.roc_auc_score(target_test, probabilities_test)
          print(auc_score)
```

0.6918361593980118

```
In [106]: fpr, tpr, _ = metrics.roc_curve(target_test, probabilities_test)
          plt.plot(fpr,tpr)
          plt.xlabel('False Positive Rate')
          plt.ylabel('True Positive Rate')
          plt.title('2-class ROC curve: AUC={0:0.2f}'.format(
                    auc_score))
          plt.show()
```



Overall our ROC curve indicates our model performs better than a naive guess, with an AUC=0.69. However, AUC is not always the best metric for unbalanced class cases such as this due to the inherent imbalance between false positives and true negatives in such a scenario.

We will now turn to precision-recall curves. Precision-recall curves are valuable because, especially in the unbalanced class case, they give the practitioner a better sense of how well the model is representing the underpresented class.


```
In [107]: average_precision = metrics.average_precision_score(target_test, probabilities_test)
          print(average_precision)
```

```
0.3509579932267077
```

```
In [108]: precision, recall, _ = metrics.precision_recall_curve(target_test, probabilities_test)
```

```
In [109]: step_kwargs = ({'step': 'post'}
                        if 'step' in signature(plt.fill_between).parameters
                        else {})
          plt.step(recall, precision, color='b', alpha=0.2,
                  where='post')
          plt.fill_between(recall, precision, alpha=0.2, color='b', **step_kwargs)

          plt.xlabel('Recall')
          plt.ylabel('Precision')
          plt.ylim([0.0, 1.05])
          plt.xlim([0.0, 1.0])
          plt.title('2-class Precision-Recall curve: AP={0:0.2f}'.format(
                  average_precision))
          plt.show()
```



Here we find average precision over all possible decision thresholds to be .35. Overall this is lower than we would like, indicating a lot of false positives are being generated from our model. However, given the business case this may be more desirable than decreased recall and increased precision as was seen when tweaking our hyperparameters.

3.0.7 Final Model Comments

Overall our model did not seem to perform as well as hoped, even with hyper-parameter tuning and feature selection. There are several next steps that could be taken to improve the model, including:

- 1) Trying different classifiers aside from logistic regressions.
- 2) Bringing in new features or datasets to improve precision/recall.
- 3) Further exploring feature engineering in the given dataset.
- 4) Gathering more data.

All of these steps could be used to help improve accuracy, precision, and recall rates of our model.